

We wish to approximate an integral of the form

$$\mathcal{F}(\nu) = \int_0^\infty dt e^{i\nu t} f(t) \quad (1)$$

for a function $f(t)$ decays to zero over a timescale τ . Taking N samples $f_n = f(t_n)$ at equally spaced times $t_n = 0, T/N, 2T/N, \dots, (N-1)T/N$ where $T > \tau$,

$$\mathcal{F}(\nu) \approx \int_0^T dt e^{i\nu t} f(t) \quad (2)$$

$$\approx \sum_{n=0}^{N-1} e^{i\nu t_n} f_n \Delta t \quad (\Delta t = T/N) \quad (3)$$

$$= \Delta t \sum_{n=0}^{N-1} e^{i\nu n T/N} f_n \quad (4)$$

This should be compared to the definition of the FFT, as [implemented in numpy](#):

$$F_k := \sum_{n=0}^{N-1} f_n e^{-2\pi i n k / N}, \quad k = 0, \dots, N-1 \quad (5)$$

We see that

$$\mathcal{F}(-2\pi k/T) = F_k \cdot \Delta t, \quad k = 0, \dots, N-1 \quad (6)$$

i.e. evaluating $F_k \cdot \Delta t$ (F_k the k^{th} Fourier component) gives the required function, but at frequencies

$$\nu = -\frac{2\pi}{T} k = -\frac{2\pi}{N\Delta t} k \quad (7)$$

By periodicity under $k \rightarrow k + N$, k can range over any N consecutive integers, however by default `np.fft.fft` provides the output in a particular ‘standard’ order. Referring to a range of N integers centered on zero, this is: ‘zero frequency, positive frequencies, negative frequencies’ i.e. $k = 0, k = 1, 2, \dots, (N-1)/2, k = -(N-1)/2, \dots, -1$ for N odd.¹ Applying `np.fft.fftshift()` to the result of `np.fft.fft` moves the elements into the more sane order of strictly increasing frequencies.

You can construct the frequencies directly using (7) or via `np.fft.fftshift(np.fft.fftfreq(N, d=dt))` which returns the frequencies (ascending order) with spacing $1/(N\Delta t)$ so that only multiplication by -2π is required. Note the maximum frequency scales as $1/\Delta t$ i.e. higher resolution in real space results in a larger range of frequencies in reciprocal space,² whilst the frequency resolution scales as $1/T$ i.e. is set by the sample time.

Using ifft To avoid the complication of having to negate the frequencies, it is simpler to use the inverse Fourier transform provided by numpy with the option `norm='forward'`, which gives

$$F_k = \sum_{n=0}^{N-1} f_n e^{2\pi i n k / N} \quad (8)$$

so that our frequencies are in the right order. **In summary, use the code**

¹For N even it goes $k = 0, k = 1, 2, \dots, N/2-1, k = -N/2, -(N-2)/2, \dots, -1$ i.e. the negative range gets the $-N/2$ value (Nyquist frequency). A subsequent `fftshift` would leave you with $k = -N/2, -N/2+1, \dots, N/2-1$.

²Although this does not necessarily provide additional information; beyond a certain frequency (determined by the most rapidly changing part of f) the spectrum will be zero.

```
# times - array of equally spaced times at which function was sampled
dt = times[1] - times[0]
N = len(sample) # for higher frequency resolution, consider oversampling
sample_fft = np.fft.fftshift(np.fft.ifft(sample, norm='forward', n=N))
nus = 2 * np.pi * np.fft.fftshift(np.fft.fftfreq(N), d=dt)
```

Note the asymmetry of the frequency spectrum when N is even (see footnote 1) mean the output here will start from $\nu = -(2\pi/T)(N/2)$ compared to $\nu = -(2\pi/T)((N-1)/2)$ if one was using the corresponding call to `fft`:

```
sample_fft = np.flip(np.fft.fftshift(np.fft.fft(sample, n=N)))
nus = np.flip(- 2 * np.pi * np.fft.fftshift(np.fft.fftfreq(N), d=dt))
```

by virtue of `np.flip`, which was required to put the negated frequencies `nus` into ascending order and move the corresponding elements of `sample_fft` in the same way.

Zero-padding For a function known to have decayed to 0 by the final sample time i.e. $f(t) = 0 \forall t \geq T$, you may consider appending zeros to the array $\{f_n\}_n^N$ of samples in order to increase the resolution in frequency. As commented above, the `n` parameter of `np.fft.ifft` can be used for this purpose: setting $n = M$ for $M > N$ results in the sample being padded with $M - N$ zeros (for $M < N$ it is cropped). Coincidentally, the FFT algorithm is most efficient when the number of samples is a power of 2, so you may want to consider setting

```
M=2**((int(np.ceil(np.log2( FAC * N ))))
```

where `FAC` is the desired oversample factor.

Endpoint correction Computing integrals in this way is susceptible to truncation error, in particular at large positive or negative ν . This is discussed in Ref. [1, Sec. 13.9]. An endpoint correction scheme based on this text is implemented in `improved_fft.py`.

References

- [1] W. Press, S. Teukolsky, W. Vetterling and B. Flannery, *Numerical Recipes: The Art of Scientific Computing*, 3e. Cambridge University Press (Cambridge, UK, 2007).